

COMP 3958 Assignment

Instructions:

- Submit a zip file containing the 5 source files `graph.ml`, `graph.mli`, `util.ml`, `util.mli` and `main.ml`.
- To focus on the functional features of OCaml, **you are not allowed to use any of its imperative or object-oriented features in this assignment. In particular, this means you cannot use references, arrays, records with mutable fields, for-loops or while-loops.** If your submission does not meet this restriction, you will get little or no credit for this assignment.
- Your files must build with no errors using `ocamlbuild` (see last paragraph); otherwise, you may receive no credit for the assignment as there may be no way to test your program.
- Provide comments to clarify your code. Additional instructions will be announced if necessary.

In this assignment, you are asked to implement Prim's algorithm for finding the minimum spanning tree of a connected, weighted and undirected graph.

For our purpose, a graph consists of vertices (labelled by strings) and edges with each edge having a weight (a positive integer) and linking two vertices. Hence there are two strings and an integer associated with an edge: the two vertices it is linking and its weight. In this write-up, all graphs are weighted and undirected, and we will use $\langle v1, v2, w \rangle$ to denote an edge linking $v1$ and $v2$ with weight w . Note that $\langle v1, v2, w \rangle$ and $\langle v2, v1, w \rangle$ denote the same edge since edges are undirected.

You are asked first to implement an abstract data type in a module named `Graph` to represent weighted undirected graphs using the "adjacency lists" representation. For this, you must use a map (from the `Map` module) to map each vertex to the list of its neighbouring vertices and their weights. For example, adding the edge $\langle "A", "B", 1 \rangle$ to the graph will add the tuple $(\text{"B"}, 1)$ to the list associated with the key `"A"` in the map, and since edges are undirected, it will also add the tuple $(\text{"A"}, 1)$ to the list associated with the key `"B"`. Note that since the module is named `Graph`, its implementation must be in a file named `graph.ml` and its signatures in a file named `graph.mli`. The file `graph.mli` is provided with this assignment — *do not modify it*.

The signature of `Graph` specified in `graph.mli` is basically as follows:

```
type t                                (* the abstract graph type *)
type edge = string * string * int     (* the edge type *)
exception Error of string              (* exception raised by some of the functions *)

val empty : t                         (* the empty graph *)

(* adds an edge to a graph
 * raises an exception if edge is invalid (one or both of the vertices
 *   are empty strings, or the weight is not positive)
 * raises an exception when a conflicting edge is being added e.g.
 *   empty |> add_edge ("A", "B", 1) |> add_edge ("B", "A", 2)
 * or when an edge joining a vertex to itself is being added
 * a "no-op" when the "same" edge is added a second time
 *)
val add_edge : edge -> t -> t

(* returns a graph from the specified list of edges; may raise an exception *)
val of_edges : edge list -> t

(* returns the list of all vertices in a graph sorted in ascending order *)
val vertices : t -> string list

(* returns whether a string is the name of a vertex in a graph *)
val is_vertex : string -> t -> bool

(* returns a sorted list of all edges in a graph; see write-up for example *)
val edges : t -> edge list

(* returns sorted list of the neighbours of a specific vertex in a graph;
 * each neighbour is represented by a (vertex, weight) pair and the list
 * is sorted in ascending order of the neighbour vertices;
 * returns an empty list if the given vertex is not in the graph *)
val neighbours : string -> t -> (string * int) list
```

Note that not all functions in `Graph` are used in Prim's algorithm.

Prim's algorithm is used to find the minimum spanning tree given a starting vertex and a graph. Refer to Wikipedia for a description of the algorithm:

https://en.wikipedia.org/wiki/Prim%27s_algorithm

The basic idea is as follows:

Input: connected undirected weighted graph `G` and a starting vertex `s`

Output: edges of a minimum spanning tree of `G`

`X := {s}` # singleton set

`T := {}` # spanning tree; empty at first

while there is an edge `<u, v, w>` with `u` in `X`, `v` not in `X` do

`<u', v', w'> :=` such an edge with minimum weight `w'`

 add `v'` to `X` # note: `u'` is already in `X`

 add edge `<u', v', w'>` to `T`

return `T`

Note that the `Set` module may be useful in your implementation.

Put your implementation of Prim's algorithm in `util.ml`. The corresponding `util.mli` file specifies 2 functions:

```
val min_tree : string -> Graph.t -> Graph.edge list * int
val read_data : string -> Graph.t
```

The `min_tree` function returns (the list of edges in) a minimum spanning tree together with its total weight. It takes the starting vertex and a graph as its arguments. It (or a helper function) implements Prim's algorithm. If there is no minimum spanning tree (e.g., the starting vertex is not in the graph), it returns `([], 0)`.

The `read_data` function is used to read information about a graph from a file. It takes the name of a file and returns an abstract graph. (It could raise the `Graph.Error` exception.) It expects each line of the input file to contain at least three words, the third of which is a positive integer. If a line does not satisfy this condition, it is skipped. For example, the content could be

```
A B 1
A C 4 # the following line is invalid and is skipped
A D -4
A D 3
B D 2
C D 5 # the following line is also invalid & is skipped
E F
```

In the above, the first line represents an edge linking vertex "A" with vertex "B" with weight 1. We can similarly interpret the other valid lines. Extra words on each line are regarded as part of a comment.

Assuming that `g` is the graph returned by `Util.read_data` from the above file content and assuming we have loaded the relevant modules, then

- `Graph.vertices g` returns `["A"; "B"; "C"; "D"]`
Note sorted order.
- `Graph.neighbours "A" g` could return `[("B", 1); ("C", 4); ("D", 3)]`
Note sorted order.
- `Graph.edges g` returns
`[("A", "B", 1); ("A", "C", 4); ("A", "D", 3); ("B", "A", 1); ("B", "D", 2);`
`("C", "A", 4); ("C", "D", 5); ("D", "A", 3); ("D", "B", 2); ("D", "C", 5)]`

The returned list of edges is in ascending order of their first vertices. For edges with the same first vertex, they are then in ascending order of their second vertices.

- `Util.min_tree "A" g` could return `[("A", "C", 4); ("B", "D", 2); ("A", "B", 1)], 7)`
Note that the minimum spanning tree may not be unique. Note also that since edges are undirected, `("A", "C", 4)` could also be `("C", "A", 4)`, etc.

You will also need to implement a main program (in `main.ml`). It checks for a starting vertex followed by a filename on the command-line, reads the file to create a graph object, calls `Util.min_tree` with the specified starting vertex and graph, and prints the result. Sample output for the above file content:

```
<A, C, 4> <B, D, 2> <A, B, 1> (weight: 7)
```

We will be using `ocamlbuild` to build your program: the main program using `ocamlbuild main.native`, the 2 other modules using `ocamlbuild graph.cmo` and `ocamlbuild util.cmo`. Make sure these commands work without warnings or errors.